

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
df=pd.read_csv("/Users/dhanashreesunilyadav/Desktop/Book1.csv")
from sklearn.model_selection import train_test_split
print(df)
print(df.head())
print(df.tail())
```

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30
5	1.5	20
6	9.2	88
7	5.5	60
8	8.3	81
9	2.7	25
10	7.7	85
11	5.9	62
12	4.5	41
13	3.3	42
14	1.1	17
15	8.9	95
16	2.5	30
17	1.9	24
18	6.1	67
19	7.4	69
20	2.7	30
21	4.8	54
22	3.8	35
23	6.9	76
24	7.8	86

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30

	Hours	Scores
20	2.7	30
21	4.8	54
22	3.8	35
23	6.9	76
24	7.8	86

In [2]:

```
print(df.info())  
print(df.shape)
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 25 entries, 0 to 24  
Data columns (total 2 columns):  
 #   Column  Non-Null Count  Dtype  
---  ---  
 0   Hours   25 non-null      float64  
 1   Scores  25 non-null      int64  
dtypes: float64(1), int64(1)  
memory usage: 528.0 bytes  
None  
(25, 2)
```

In [3]:

```
print(df.groupby("Hours").size())
```

```
Hours  
1.1    1  
1.5    1  
1.9    1  
2.5    2  
2.7    2  
3.2    1  
3.3    1  
3.5    1  
3.8    1  
4.5    1  
4.8    1  
5.1    1  
5.5    1  
5.9    1  
6.1    1  
6.9    1  
7.4    1  
7.7    1  
7.8    1  
8.3    1  
8.5    1  
8.9    1  
9.2    1  
dtype: int64
```

In [4]:

```
print(df.groupby("Scores").size())
```

Scores

```
17    1
20    1
21    1
24    1
25    1
27    1
30    3
35    1
41    1
42    1
47    1
54    1
60    1
62    1
67    1
69    1
75    1
76    1
81    1
85    1
86    1
88    1
95    1
dtype: int64
```

In [5]:

```
print((df.isnull()).sum())
```

```
Hours    0
Scores   0
dtype: int64
```

In [6]:

```
print(df.describe())
```

	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000

In [7]:


```
print(df.corr())
```

	Hours	Scores
Hours	1.000000	0.976191
Scores	0.976191	1.000000

In [8]:

```
print(df.plot(x="Hours",y="Scores",style="o"))
```

AxesSubplot(0.125,0.125;0.775x0.755)



In [9]:

```
x=df.iloc[:, :-1].values  
print(x)
```

```
[[2.5]  
[5.1]  
[3.2]  
[8.5]  
[3.5]  
[1.5]  
[9.2]  
[5.5]  
[8.3]  
[2.7]  
[7.7]  
[5.9]  
[4.5]  
[3.3]  
[1.1]  
[8.9]  
[2.5]  
[1.9]  
[6.1]  
[7.4]  
[2.7]  
[4.8]  
[3.8]  
[6.9]  
[7.8]]
```

In [10]:

```
y=df.iloc[:,1].values  
print(y)
```

```
[21 47 27 75 30 20 88 60 81 25 85 62 41 42 17 95 30 24 67 69 30 54 35  
76  
86]
```

In [11]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)
```

In [12]:

```
from sklearn.linear_model import LinearRegression  
regressor=LinearRegression()  
regressor.fit(x_train,y_train)
```


Out[12]:

```
LinearRegression()
```

In [13]:


```
print(plt.scatter(x,y))
```

<matplotlib.collections.PathCollection object at 0x7f879d94ac70>



In [14]:

```
line=regressor.coef_*x+regressor.intercept_  
  
plt.scatter(x,y)  
plt.plot(x,line);
```



In [15]:

```
print(x_test)  
y_pred=regressor.predict(x_test)  
#print(y_pred)
```

```
[[1.5]  
 [3.2]  
 [7.4]  
 [2.5]  
 [5.9]  
 [3.8]  
 [1.9]  
 [7.8]]
```

In [16]:

```
dataframe=pd.DataFrame({'Actual':y_test,'Predicted':y_pred})  
print(dataframe)
```

	Actual	Predicted
0	20	17.053665
1	27	33.694229
2	69	74.806209
3	30	26.842232
4	62	60.123359
5	35	39.567369
6	24	20.969092
7	86	78.721636

In [17]:

```
study_hour=[[9.25]]
score_prediction=regressor.predict(study_hour)

print('Number of Hours of studying:',study_hour)
print('Predicted Score:',score_prediction)
```

Number of Hours of studying: [[9.25]]
Predicted Score: [92.91505723]

In [18]:

```
from sklearn.metrics import r2_score
print("R2 Score:", r2_score(y_test,y_pred))
```

R2 Score: 0.9568211104435257

In [19]:

```
from sklearn.metrics import mean_squared_error
print('Mean Square Error:', mean_squared_error(y_test,y_pred))
```

Mean Square Error: 22.965097212700428

In [20]:

```
from sklearn import metrics
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test,y_pred))
```

Mean Absolute Error: 4.419727808027651

In [21]:

```
#If a student studies 9.25 hrs per day then he would score 92.91
```